

STAT 513/413: Lecture 12

Classical Monte Carlo

(probabilistic theorems)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Recall: classical Monte Carlo

Classical Monte Carlo deals with the estimation of expected values, with probabilities also falling under this case. This particular problem allows for more than merely intuitive approach: it is possible to assess the precision more rigorously via existing mathematical theorems.

Notation: we are still after quantity ϑ
we still compute its estimate T_N

And the subscript N still indicates how many random numbers were used in the computation of T_N : the higher N , the more numbers used

But now we have $\vartheta = E(T_N)$

Or even more precisely: we are after $\vartheta = E(X)$ where X has the same distribution as random numbers X_1, X_2, \dots, X_N ; that is

$$\vartheta = E(X_1) = E(X_2) = \dots = E(X_N)$$

This particular case covers quite a lot of instances; in particular...

Probabilities come as a special case of this scheme

We are interested in $p = P(E)$, the probability of some event E

We define an *indicator random variable* Y to be

$$Y = \begin{cases} 1 & \text{if event } E \text{ happens} \\ 0 & \text{if event } E \text{ does not happen} \end{cases}$$

The expected value of Y is

$$E(Y) = 1P[Y = 1] + 0P[Y = 0] = P[Y = 1] = p$$

Clear?

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

We estimate expected values by averages

(or means by means - but this formulation may confuse some)

We make (that is, generate) N replications of X

$$X_1, X_2, \dots, X_N$$

are considered independent and having the same distribution as X

Then

Assignment Project Exam Help

$$\vartheta = E(X_1) = E(X_2) = \dots = E(X_N)$$

<https://powcoder.com>

and

$$T_N = \bar{X} = \frac{1}{N} \sum_{i=1}^N X_i$$

Add WeChat powcoder

Note that then T_N is an unbiased estimator of ϑ

$$E(T_N) = E(\bar{X}) = \frac{1}{N} \sum_{i=1}^n E(X_i) = \frac{1}{N} (N\vartheta) = \vartheta$$

Why does Monte Carlo work in this case? A classical result:

Jacob Bernoulli (1655-1705), *Ars Conjectandi* (1713)

Law of large numbers (in the version proved later by Kolmogorov)

If X_1, X_2, \dots, X_N are independent random variables with the same distribution, such that the expected value of all of them (or of one: same distribution!) [exists and] is finite

$$\mu = E(X_1) = E(X_2) = \dots$$

then the average $\bar{X} = \frac{1}{N} \sum_{i=1}^N X_i$

<https://powcoder.com>

converges to the mean μ for $n \rightarrow \infty$, *in probability*:

for every $\varepsilon > 0$, $P[|\bar{X} - \mu| \geq \varepsilon] \rightarrow 0$ with $N \rightarrow \infty$

If our estimated value is $\vartheta = \mu$ and its estimate is $T_N = \bar{X}$, then this establishes consistency, right?

(A fine point: mere T_N being unbiased could not be enough, we still need something more... and there are also another fine points for math enthusiasts)

For math enthusiasts

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Probabilities are again a special case

We are interested in $p = P(E)$, the probability of some event E

We define *indicator random variables* to be

$$Y_i = \begin{cases} 1 & \text{if event } E \text{ happens in the } i\text{-th repetition (trial)} \\ 0 & \text{if event } E \text{ does not happen in the } i\text{-th repetition} \end{cases}$$

If the repetitions are independent, then also Y_i are independent and their common mean is

$$E(Y_i) = 1P[Y_i = 1] + 0P[Y_i = 0] = P[Y_i = 1] = p$$

So, then we calculate

$$\bar{Y} = \frac{1}{N} \sum_{i=1}^N Y_i = \frac{\text{card}\{i : Y_i = 1\}}{N}$$

to estimate (that is, compute) p

The consistency follows from the Law of large numbers

Let us try it on toy examples

Recall: what is the probability of obtaining 2 or less heads in 3 coin tosses? (Answer: 0.875)

A problem for expected values proper (not probabilities): what is the expected value of heads in 3 coin tosses?

(Well, everybody knows that 1.5 - but we need a simple problem to start with.)

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

First: it works for estimating a probability

(We should have kept the seed. Just in case.)

```
> mean(apply(replicate(1000,rbinom(3,1,0.5)),2,sum) <= 2)
[1] 0.858
```

```
> mean(apply(replicate(1000,rbinom(3,1,0.5)),2,sum) <= 2)
[1] 0.854
```

```
> mean(apply(replicate(1000,rbinom(3,1,0.5)),2,sum) <= 2)
[1] 0.888
```

```
> mean(apply(replicate(10000,rbinom(3,1,0.5)),2,sum) <= 2)
[1] 0.8806
```

```
> mean(apply(replicate(10000,rbinom(3,1,0.5)),2,sum) <= 2)
[1] 0.8761
```

```
> mean(apply(replicate(1000000,rbinom(3,1,0.5)),2,sum) <= 2)
[1] 0.875105
```

```
> mean(apply(replicate(1000000,rbinom(3,1,0.5)),2,sum) <= 2)
[1] 0.874815
```

```
> mean(apply(replicate(1000000,rbinom(3,1,0.5)),2,sum) <= 2)
[1] 0.875326
```

```
> mean(apply(replicate(10000000,rbinom(3,1,0.5)),2,sum) <= 2)
[1] 0.8748622
```

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

Works for the expected values as well

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder

The evaluation of precision

Classical Monte Carlo has one big advantage over the intuitive approach: we can evaluate the precision in an exact way

For starters, let us - *tentatively!* - measure it once again by the mean squared error

$$E[(T_N - \vartheta)^2]$$

(Why? Because it is quite reasonable and we are able to handle the mathematics - compared to other possible choices.)

An important case is when $E(T_N) = \vartheta$ for all ϑ under consideration

Recall: such estimators are called *unbiased*

For unbiased estimator, the mean squared error is equal to their

variance
$$E[(T_N - \vartheta)^2] = E[(T_N - E(T_N))^2] = \text{Var}(T_N)$$

Sometimes, we may also look at the square root of variance

standard deviation
$$\sqrt{\text{Var}(T_N)}$$

Performance guarantees II: precision via variance

Variance has direct consequences for the precision via the **Chebyshev(-Bienaymé) inequality**, which in its general form says that for any random variable Z with both $E(Z)$ and $\text{Var}(Z)$ finite

$$P[|Z - E(Z)| \geq \varepsilon] \leq \frac{\text{Var}(Z)}{\varepsilon^2}$$

Assignment Project Exam Help

Thus, if T_N is an unbiased estimator of ϑ , then for every $\varepsilon > 0$,

$$P[|T_N - \vartheta| \geq \varepsilon] \leq \frac{\text{Var}(T_N)}{\varepsilon^2}$$

<https://powcoder.com>

Add WeChat powcoder

So, the variance (or its square root, standard deviation) is an important indicator of the precision of the estimate T_N of ϑ

And it is evaluated *explicitly* via Chebyshev inequality, so *no mean square error anymore!* We can do better here!

A side product: if T_N is unbiased and we can show mathematically that $\text{Var}(T_N) \rightarrow 0$ for $N \rightarrow \infty$, then consistency for T_N immediately follows

So, again the canonical case

Let X_1, X_2, \dots, X_N be random variables

with the same distribution

they (all) have (the same) expected value μ

and they also have the same variance σ^2

Assignment Project Exam Help

What is the variance of $\bar{X} = \frac{1}{N} \sum_{i=1}^N X_i$? Standard calculation...

<https://powcoder.com>

$$\text{Var}(\bar{X}) = \text{Var}\left(\frac{1}{N} \sum_{i=1}^N X_i\right) = \frac{1}{N^2} \sum_{i=1}^N \text{Var}(X_i) = \frac{1}{N^2} N \sigma^2 = \frac{\sigma^2}{N}$$

This may serve as the measure of precision

- only we need to know σ^2

(which we may not, however)

Performance guarantees II: \bar{X} via Chebyshev

Once we know σ^2 , we have a bound on the probability that our results deviates from μ by more than ε :

Knowing that $E(\bar{X}) = \vartheta$, we apply the Chebyshev inequality to \bar{X}

$$P[|\bar{X} - \vartheta| \geq \varepsilon] \leq \frac{\text{Var}(\bar{X})}{\varepsilon^2} = \frac{\sigma^2}{N\varepsilon^2}$$

This gives us a lucid expression of precision (as long as we know σ^2)

Also, note that, as $\frac{\sigma^2}{N\varepsilon^2} \rightarrow 0$ with $N \rightarrow \infty$ (whatever σ^2 is)

we have automatically consistency here

Return to precision: what if σ^2 is not known (as is often the case)?

An upper bound can help

First, recall “the other formula” for the variance: this yields

$$\text{Var}(Z) = E(Z^2) - (E(Z))^2 \leq E(Z^2)$$

Note: as $E(Z)$ is usually finite (we are estimating $E(Z)$, so if it would not be properly defined, it would be hopeless), the crux is in $E(Z^2)$

And geeks know: $E(Z^2) < \infty$ implies not only $\text{Var}(Z) < \infty$, but also that $E(|Z|) < \infty$, that is, $E(Z)$ exists and is finite

Recall some very simple, but also very useful principles: if for random variables U and V we have $U \leq V$ (with probability one)

$$\text{then } E(U) \leq E(V)$$

In particular, if $X^2 \leq Y^2$ (with probability one)

$$\text{then } E(X^2) \leq E(Y^2)$$

(Beware: nothing like that necessarily follows for variances)

A very often used upper bound on variance

If a random variable Z has values in $[a, b]$ (with probability one)

then
$$\text{Var}(Z) \leq \frac{(b - a)^2}{4} \quad (\text{try to prove that!})$$

Combined with the Chebyshev inequality for \bar{X} , we get

Assignment Project Exam Help

$$P[|\bar{X} - \vartheta| \geq \varepsilon] \leq \frac{\text{Var}(\bar{X})}{\varepsilon^2} \leq \frac{(b - a)^2}{4n\varepsilon^2}$$

<https://powcoder.com>

Add WeChat powcoder
whenever $\vartheta = E(\bar{X})$ and all X_i forming \bar{X} have values in $[a, b]$
(with probability one)

Probabilities are a special case

In particular, for probabilities:

we define Y_i with values 0 and 1 as above

these Y_i happen to have expected value equal to p and

$$\text{Var}(Y_i) = (1-p)(0-p)^2 + p(1-p)^2 = p(1-p)^2 - p^2(1-p) = p(1-p)$$

Assignment Project Exam Help

Given that p is in $[0, 1]$, we have an inequality $p(1-p) \leq 1/4$

<https://powcoder.com>

$$\text{And therefore } \text{Var}(\bar{Y}) = \frac{p(1-p)}{N} \leq \frac{1}{4N}$$

Add WeChat powcoder

So, for 0-1 Y_i , we get now from the Chebyshev inequality

$$P[|\bar{Y} - \mu| \geq \varepsilon] \leq \frac{\sigma^2}{N\varepsilon^2} \leq \frac{1}{4N\varepsilon^2}$$

Return to precision: how is this used?

We first select ε to express precision we would like to achieve
for instance, if we want 2 decimal digits, then $\varepsilon = 0.005$
(or slightly less)

Then, we select “margin of error” δ : if we are to be outside the bounds with some probability, what probability we would tolerate
for instance, $\delta = 0.01$

We are interested in N that will guarantee us chosen ε and δ

The Chebyshev inequality says that it will be equal to or larger than

N satisfying $\frac{\sigma^2}{N\varepsilon^2} = \delta$ which is $N = \frac{\sigma^2}{\delta\varepsilon^2}$

if estimating a probability we need $N \geq \frac{1}{4\delta\varepsilon^2} = 1000000$

Note: to increase precision by one decimal digit

we have to divide ε by 10, hence multiply N by 100

and to decrease “margin of error” by one decimal digit

we have to divide δ by 10, hence multiply N by 100

Performance guarantees II: \bar{X} for bounded X_i

Hoeffding inequality. Suppose that X_1, X_2, \dots, X_N are independent random variables, each X_i bounded by an interval $[a_i, b_i]$ (that is, X_i takes values in $[a_i, b_i]$ with probability one).

If $\vartheta = E(\bar{X})$, then

$$P[|\bar{X} - \vartheta| \geq \varepsilon] \leq 2e^{-\frac{2N\varepsilon^2}{\sum_{i=1}^N (b_i - a_i)^2}}$$

In particular, if all $[a_i, b_i] = [a, b]$, then

$$P[|\bar{X} - \vartheta| \geq \varepsilon] \leq 2e^{-\frac{2N\varepsilon^2}{(b-a)^2}}$$

And then, if $[a, b] = [0, 1]$,

$$P[|\bar{X} - \vartheta| \geq \varepsilon] \leq 2e^{-2N\varepsilon^2}$$

(the latter case is also known as the **Chernoff bound**)

The use of this one

Suppose that we have $[a_i, b_i] = [a, b]$ - and as before

we want precision ε and the “margin of error” δ

To have N guaranteeing this, it has to be equal or greater than

the one satisfying
$$N = -\frac{(b-a)^2}{2\varepsilon^2} \log\left(\frac{\delta}{2}\right)$$

for indicator variables it is
$$N = -\frac{1}{2\varepsilon^2} \log\left(\frac{\delta}{2}\right)$$

So, for $\delta = 0.01$ and $\varepsilon = 0.005$ we get $N \geq 105967$

Note: to increase precision by one decimal digit

we have to divide ε by 10, hence multiply N by 100

this is the same of the Chebyshev

But, decreasing “margin of error” by one decimal digit:

dividing δ by 10 means *adding* $\frac{\log 10}{2\varepsilon^2} \approx \frac{1.1513}{\varepsilon^2}$ to N

That looks cheaper...

Some remarks

Note first:

All the mentioned inequalities bound some *probability* from above
thus, to be nontrivial, the bound has to be < 1

Some examples now:

For $\delta = 0.01$ and $\varepsilon = 0.005$ we get $N \geq 105967$ by the Hoeffding

In the same situation, for $N = 1000000$ and $\varepsilon = 0.005$, Hoeffding
says we $\delta \leq 3.8575 \times 10^{-22}$

Great, this is way much we hoped for, but nonetheless fine... Can
we expect also better precision here - maybe with slightly higher
probability of error?

If we try $\varepsilon = 0.0005$, still with $N = 1000000$, we actually obtain the
bound $\delta \leq 1.213$ - which does not say anything nontrivial: it is a
bound on a probability which is always < 1

To achieve $\delta = 0.01$ for $\varepsilon = 0.0005$, we need $N \geq 10596634$

And also: Hoeffding works for *bounded* X_i

but in that case we also have an upper bound on variance!

Chebyshev strikes back

If X_i is within $[a, b]$ (with probability one)

then Chebyshev and the inequality on variance together yield

$$P[|\bar{X} - \vartheta| \geq \varepsilon] \leq \frac{\text{Var}(\bar{X})}{\varepsilon^2} \leq \frac{(b - a)^2}{4N\varepsilon^2}$$

and Hoeffding inequality gives

$$P[|\bar{X} - \vartheta| \geq \varepsilon] \leq 2e^{-\frac{2N\varepsilon^2}{(b-a)^2}} = e^{-\frac{1}{2} \frac{4N\varepsilon^2}{(b-a)^2}}$$

So, the bounds are respectively $\delta = \frac{(b-a)^2}{4N\varepsilon^2}$ and $2e^{-\frac{1}{2\delta}}$


Analyzing this, we figure out that for small δ , the second bound is better; the situation reverses at about $\delta \approx 0.2322026$

However, we are interested in *small* δ

and in such situation is Hoeffding better

Finally: for given precision ε , we decrease the “margin of error” by increasing N - but we may also do it by decreasing $(b - a)^2$ or σ^2

A pictorial comparison of bounds



The comparison of the required N's

For a given precision ε , the required N is greater or equal than

$$\frac{(b-a)^2}{4\delta\varepsilon^2} = \frac{(b-a)^2}{2\varepsilon^2} \left(\frac{1}{2\delta} \right) \quad \text{by the Chebyshev and bound}$$

$$-\frac{(b-a)^2}{2\varepsilon^2} \log \left(\frac{\delta}{2} \right) \quad \text{by the Hoeffding}$$

Once ε is given, then $(b-a)^2/(2\varepsilon^2)$ is given as well

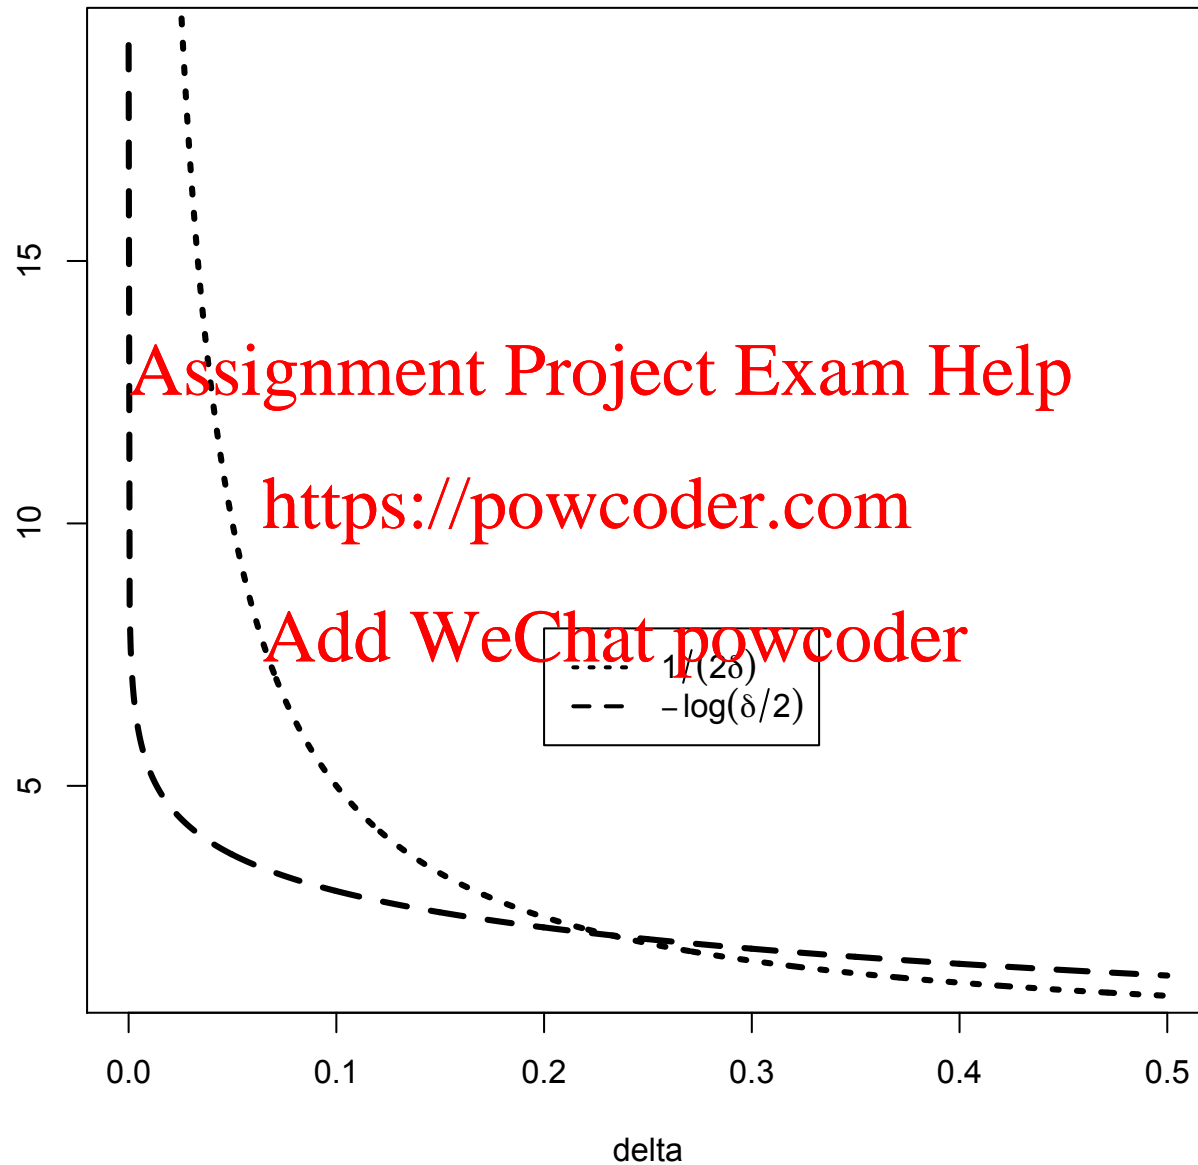
and we can compare $\frac{1}{2\delta}$ and $-\log \left(\frac{\delta}{2} \right)$

Again, they are equal for $\delta \approx 0.2322026$

and then for smaller δ Hoeffding gives smaller N

in fact, *much* smaller for very small δ

The comparison of required N for probabilities



Sometimes, however, σ^2 can only be estimated

The well-known estimate of σ^2 is $s^2 = \frac{1}{N-1} \sum_{i=1}^n (X_i - \bar{X})^2$

and $\text{Var } \bar{X}$ is then estimated by $\frac{s^2}{N} = \frac{\frac{1}{N-1} \sum_{i=1}^n (X_i - \bar{X})^2}{N}$

Assignment Project Exam Help

Note: sometimes instead of $s^2 = \frac{1}{N-1} \sum_{i=1}^n (X_i - \bar{X})^2$

Add WeChat powcoder

they use $\frac{1}{N} \sum_{i=1}^n (X_i - \bar{X})^2$

(compare textbook, pp. 151-152 and also pp. 52-53)

The reason is that - in the present context - N is usually large...

...so $\frac{1}{N}$ differs very little from $\frac{1}{N-1}$

In fact, it may work that way for probabilities

Suppose again that Y_1, Y_2, \dots, Y_N are independent random variables
such that $P[Y_1 = 1] = p$ and $P[Y_1 = 0] = 1 - p$

We know that $\vartheta = p = E(Y_i)$

and also that $\sigma^2 = \text{Var}(Y_i) = p(1 - p)$

We estimate ϑ by $\bar{Y} = \frac{1}{N} \sum_{i=1}^N Y_i$

<https://powcoder.com>

So we could perhaps also estimate $\text{Var}(\bar{Y})$ by $\bar{Y}(1 - \bar{Y})$

Add WeChat powcoder

but we could do it also by $\frac{1}{N-1} \sum_{i=1}^N (Y_i - \bar{Y})^2$

and perhaps also by $\frac{1}{N} \sum_{i=1}^N (Y_i - \bar{Y})^2$

Are those related in some way?

Let us do a computational experiment first

```
> tosses = floor(2*runif(20)) ; tosses
[1] 0 1 0 1 0 0 1 0 0 1 1 0 0 1 1 0 1 1 1 0
> mean(tosses)*(1-mean(tosses))
[1] 0.25
> var(tosses)
[1] 0.2631579
> var(tosses)*19/20
[1] 0.25
```

Assignment Project Exam Help

Maybe only a coincidence... OK, another one:

<https://powcoder.com>

```
> tosses = floor(2*runif(30))
> tosses
[1] 1 1 0 0 0 1 1 0 0 0 0 0 0 0 1 1 1 0 1 0 0 0 0 0 1 1 0 0 1 1 0
> mean(tosses)*(1-mean(tosses))
[1] 0.24
> var(tosses)
[1] 0.2482759
> var(tosses)*29/30
[1] 0.24
```

Add WeChat powcoder

So, a rule here?

Your theory for that:

Assignment Project Exam Help

<https://powcoder.com>

Add WeChat powcoder